
Reinforcement Learning for Financial Time-Series Forecasting

Naman Choudhary
Carnegie Mellon University
namancho@andrew.cmu.edu

Karthik Subramaniam
Carnegie Mellon University
ksubram2@andrew.cmu.edu

Nimish Jindal
Carnegie Mellon University
nimishj@andrew.cmu.edu

Sarah Li
Carnegie Mellon University
sarahli@andrew.cmu.edu

Abstract

1 This paper presents RLFinNet, a novel reinforcement learning-based model tailored
2 for the efficient and rapid forecasting of financial time-series data. Focused on
3 the specific dynamics of individual company stock prices, RLFinNet addresses
4 the need for models that are not only accurate but also quick to adapt to new data,
5 avoiding the influence of irrelevant market factors. Our approach enhances the
6 classical deep reinforcement learning framework by incorporating an innovative
7 dual forecasting model, adapted from the OneNet architecture, which leverages
8 both cross-time and cross-variable dependencies to predict future stock prices ef-
9 fectively. We describe the adaptations made to our model to handle the multivariate
10 financial data, detailing the updates to our data loader and the modifications to the
11 training loop to accommodate sequential processing across various datasets. The
12 model was tested using a financial dataset featuring a range of metrics including
13 open, high, low, and closing prices, among others. We conducted a series of ex-
14 periments with varying prediction and sequence lengths to optimize our model
15 for different forecasting horizons. Our findings demonstrate that RLFinNet can
16 achieve significant improvements in prediction accuracy with a mean squared error
17 (MSE) goal of less than 0.4, aligning with the performance metrics of established
18 models. The results highlight the model’s potential in transforming financial market
19 analytics by providing robust, scalable, and efficient predictive insights. This work
20 not only advances the field of financial forecasting but also opens up new avenues
21 for real-time data processing in other domains.

22 1 Introduction

23 The financial portfolio management problem is the process of constant redistribution of a fund
24 into different financial products. Algorithmic trading has thus received increasing spotlight as the
25 volume of financial data grows. Algorithmic trading involves the use of mathematical models and
26 statistical analysis to exploit market opportunities and make financial trading decisions. Recently,
27 neural-network based trading has been gaining attention, as price predictions are not market actions,
28 and converting them into actions requires additional logic for real-world applications. Due to its
29 unsupervised nature, we expand upon model-free reinforcement learning for this research project.

30 Our project goal is to create an efficient and quick network – **RLFinNet** – that is small in terms of
31 architecture parameters. This is beneficial to traders as a model specialized to a specific company’s
32 stock prices will not be influenced by irrelevant companies’ stock prices, and retraining the model
33 with custom or tweaked hyper-parameters can be done in a matter of seconds. The model is ultimately

34 designed to train on real-time data and make predictions more rapidly than more saturated and
35 generalized models.

36 1.1 Model Description

37 Deep Reinforcement Learning (DLR) has been an effective candidate of time-series forecasting for
38 years. DRL as a stochastic problem can represent an environment, in which an a model “agent” is able
39 to choose a course of actions given state policies, with the ultimate goal of maximizing accumulated
40 rewards. For this project, we aim to show the usefulness of Deep Reinforcement Learning for online
41 and rapid training of time-series forecasting.

42 The base model OneNet that we have adapted is made of Convolutional layers. The entire model is
43 an ensemble of 2 model– one cross-variable and one cross-time. The individual models are trained
44 separately using MSE loss and the models are ensembled using reinforcement learning to adjust the
45 weight of the combination of the two models. This will be further elaborated in *section 5*.

46 1.2 Dataset

47 We used the **ACL18** financial dataset from the previous STOTA model, StockNet [23], which features
48 historical prices from a range of companies. Specifically, we intend to predict the future closing price
49 of a stock, based on its *Open, High, Low, Volume, Volume*, and *Adjusted Close* prices traded on the
50 current day. Including the date, the input is 7-Dimensional.

51 2 Literature Review

52 The Stock Market is a financial ecosystem involving over 100 trillion dollars in 2021, according to
53 the world bank [22]. As a result, the financial industry generates enormous data every single day.
54 Ranging from customer information to transaction records, the data are not only kept confidential but
55 also provided quickly to customers for strategic planning and decision-making. With the advancement
56 of these digital technologies, enormous datasets can be efficiently handled, analyzed, and recorded.
57 Quantitative stock investment is a fundamental financial task that highly relies on accurate prediction
58 of market status and profitable investment decision making [1]. Different investors may approach
59 the market from different angles, for example, by studying market behavior, identifying influential
60 factors, trading stocks, forecasting market directions, making asset recommendations for portfolio
61 management, etc. Despite any and all of these goals, the investors continuously face adversity with
62 time constraints.




Figure 1: Value of the global stock market [22]

63 We reference several research papers for our project. In "Enhancing Stock Movement Prediction
64 with Adversarial Training [9]," FuLi *et al* used adversarial training to improve the generalizability
65 of Neural Network predictions on stock data by training the model under intentional perturbations
66 to stochastic price variables, which had conventionally been treated as static features in prediction
67 tasks. This novel approach has demonstrated an improvement in accuracy compared to the previous
68 state-of-the-art model StockNet developed by [Xu and Cohen, 2018] to forecast temporally dependent
69 data.

70 Shaban *et al* [20] compared two novel models - One with LSTM layers and one with BiGRU layers
 71 to existing DL-based approaches for stock prediction. The new recurrent architectures proved to be
 72 computationally more efficient and proved better at real-time trading. They were able to predict stock
 73 prices 10 minutes and 30 minutes into the future. While our project is not based on real-time trading,
 74 we do believe that recurrent approaches could be much more efficient than standard time-series
 75 forecasting models.

76 Applying RL to portfolio optimization is a fairly old idea, but has never caught on in the industry.
 77 Ralph Neunier [18] employed 'Adaptive Dynamic Programming' (The lexicon of Machine Learning
 78 Engineers has changed greatly over the years) in 1995. He showed the feasibility of QL learning as a
 79 function approximator as a Markov Decision problem. Moody *et al* [16] also implemented direct
 80 reinforcement with great success. These researchers paved the way for Reinforcement Learning in
 81 Finance and Portfolio Management.

82 In 2021, Weiwei [11] summarized the recent progress of deep learning models in stock price prediction.
 83 He concluded that Graph Neural Network (GNN) and DRL models held promise in the field but were
 84 yet to be explored in depth. Cui *et al* [7] were one of the first to use convolutional neural networks
 85 (CNNs) for time-series forecasting of stock prices. It leverages the strength of CNN to automatically
 86 learn good feature representations in both time and frequency domains. In particular, MCNN contains
 87 multiple branches that perform various transformations of the time series, which extract features
 88 of different frequency and time scales, addressing the limitation of many previous works that they
 89 only extract features at a single timescale. Cheng *et al* [24] also attempted an attention-based LSTM
 90 model in stock prediction. While the results are undoubtedly impressive, it still leaves room for
 91 improvement in terms of computational efficiency.

92 Zhipeng Liang *et al* demonstrated superiority of adversarial models for portfolio management [13].
 93 They implemented three state-of-art continuous reinforcement learning algorithms, Deep Deterministic
 94 Policy Gradient (DDPG), Proximal Policy Optimization (PPO) and Policy Gradient (PG)
 95 in portfolio management. All of them are widely-used in game playing and robot control. The
 96 implemented algorithms were extremely sophisticated and took into account the volatility of the stock
 97 market and the assets (which would have influenced the investors). Chaouki *et al* [5] also explored
 98 deep deterministic learning for portfolio optimization, but also compared their approach to the known
 99 optimal strategies of the time. They concluded that their DDPG-RL approach was always able to
 100 recover the established baselines (which even investors deem difficult). It did not beat the optimal
 101 strategies but was very close to the optimal. However, one specificity of their approach was the
 102 reward function was assumed unknown. Thus, more efficient algorithms could be made that address
 103 this concern of theirs and could potentially take the model past the other 'optimal' strategies that they
 104 considered.

105 Similarly, Cannelli *et al* [4] employed Q Learning for hedging. If an agent is trained solely on
 106 simulated data, the run-time performance will primarily reflect the accuracy of the simulation, which
 107 leads to the classical problem of model choice and calibration. However, their approach was not
 108 as successful as what we saw above. Chen *et al* [6] and Liang *et al* [13] explored adversarial
 109 reinforcement learning with similar motivations as Feng *et al*.

110 In 2020, Liu *et al* [14] proposed a library of three layers: environments, agents and applications. The
 111 three layers of FinRL library are stock market environment, DRL trading agent, and stock trading
 112 applications. The agent layer interacts with the environment layer in an exploration-exploitation
 113 manner, whether to repeat prior beneficial decisions or to make new actions hoping to get greater
 114 rewards. The lower layer provides APIs for the upper layer, making the lower layer transparent to the
 115 upper layer. The environment would be a real-time time-driven trading simulator. By interacting with
 116 the environment, the trading agent will derive a trading strategy with the maximized rewards as time
 117 proceeds.

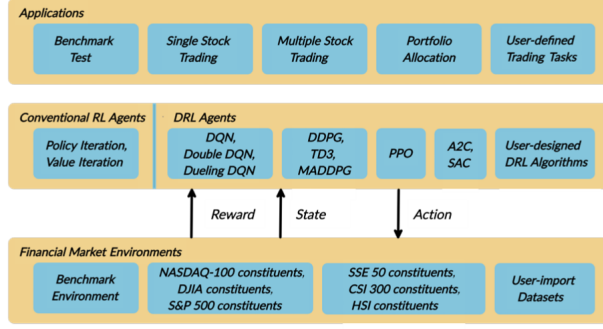


Figure 2: An overview of the FinRL Library’s 3 layers

In the realm of online time series forecasting, addressing concept drift—where the underlying patterns of data change over time—is crucial. A recent notable work in this domain is "OneNet" by Yi-Fan Zhang *et al* [24]. OneNet innovatively tackles the challenge of concept drift by dynamically updating and combining two models: one focused on temporal dependencies and the other on cross-variable dependencies. This approach leverages the strengths of both models, enabling more accurate and robust forecasting in the face of evolving data patterns. OneNet incorporates a reinforcement learning-based strategy to adjust the weights of these models dynamically, showcasing a significant improvement in forecasting accuracy, reducing mean-squared errors by over 50% compared to the STOTA methods. This method’s success underscores the potential of adaptive, multi-model strategies in handling the complexities of online forecasting, presenting a promising direction for our exploration in financial portfolio management with deep reinforcement learning.

3 Baseline Selection

In the literature, we encountered many issues with choosing a reinforcement learning model trained on finance data. While Chaouki *et al*’s theory was sound, the results were not stellar and the heuristics to evaluate the model were not standardized with the rest of the literature. Liu *et al*’s FinRL, while very informing, is more of a library of models than a novel model in itself. The implementation of these older architectures described in 2018 [14] may not be very practical as Reinforcement Learning has changed and adaptive within recent years. Azhikodan *et al* [3], Jia *et al* [10], and Meng *et al*’s [17] models also followed a similar approach using outdated reinforcement learning methods.

We used Adv-ALSTM, as suggested in [9], as the baseline implementation and run the model on two benchmarks on stock movement prediction. This adversarial learning approach helped them create a very generalized model. We were interested to see how it works since the goal of our project was to create hyper-specific models. Using a state of the art generalized model as our baseline would help validate or invalidate our hypothesis.

This baseline trained on two sets of data from gathered by StockNet [23]. The **ACL18** sources historical stock data from Jan-01-2014 to Jan-01-2016 of 88 high-trade-volume-stocks in NASDAQ and NYSE, while **KDD17** includes a longer range from Jan-01-2007 to Jan-01-2016 of 50 stocks in U.S. markets. This paper introduces adversarial training into the realm of financial market predictions, a domain characterized by high uncertainty and noise. This improves the model resilience against market volatility. The adversarial training method suggested by the paper is still referenced by many papers to this day. The paper offers a comprehensive experimental setup, including datasets, evaluation metrics, and model configurations, which can be easily used to benchmark our model’s performance. The heuristics and evaluation metrics are simple to understand and compare with other architectures, the generalization is a great advancement towards incorporating AI in portfolio management, and their approach of normalizing the stock prices to explicitly capture the interaction seems most practical.




Figure 3: Illustration of the Adversarial Attentive LSTM

4 Baseline Implementation

During our research, we referenced a baseline model described in "Enhancing Stock Movement Prediction with Adversarial Training" [9] and were able to successfully replicate its resulting accuracies as demonstrated in Table 1. For the implementation and execution of our machine learning model, we utilized Google Cloud Platform (GCP). The model itself was constructed using TensorFlow and Keras. We first ran the model without the saved checkpoint and then with the checkpoint to verify the results. The run without the saved checkpoint was able to achieve the same accuracy when compared to one with the saved checkpoint after 150 epochs. This is corroborated in Figure 4 and Figure 13. Performance over each epoch was recorded on Wandb. The accuracy at the start of pre-trained model is very close to the accuracy achieved at the end of 150 epochs of the model run without saved checkpoints.




Figure 4: Comparing model run with and without saved checkpoint for KDD17 Data




Figure 5: Comparing model run with and without saved checkpoint for KDD17 Data

Table 1: Performance comparison on the two datasets.

Method	ACL18		KDD17	
	Acc	MCC	Acc	MCC
Adv-ALSTM	57.20	0.1483	53.05	0.0523
Our implementation	58.72	0.1752	53.05	0.05292
Difference	1.52	0.0269	0	0.00062

5 Model Description

5.1 OneNet Model for Online Time Series Forecasting

In addressing the challenge of concept drift in time series forecasting, OneNet stands out for its online ensembling approach, dynamically balancing model predictions to accommodate changing patterns in data over time. Figure 2 illustrates the dual architecture of OneNet, which processes multivariate data through two specialized branches: the Cross-Time Forecaster and the Cross-Variable Forecaster, each tailored to capture temporal and cross-variable dependencies, respectively. The core innovation of OneNet lies in its Online Convex Programming (OCP) block that integrates the predictions of both forecasters. This OCP block not only draws upon the long-term historical data through Exponentiated Gradient Descent (EGD) but also adapts rapidly to recent trends via offline reinforcement learning. This dual-weighting system allows OneNet to maintain robust forecasting while dynamically adapting to concept drift, which is often seen in real-world scenarios such as financial markets.

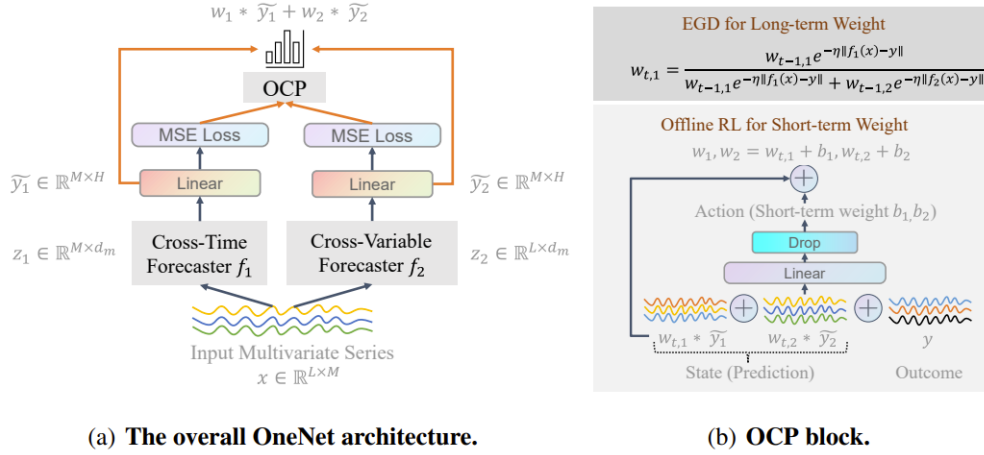


Figure 6: (a) The Overall OneNet architecture (b) OCP block [24]

Key components of OneNet Architecture:

- **Cross-Time Forecaster(f_1):** Focuses on the temporal correlation within the data, projecting the input series into a representation that captures time-dependent variations.
- **Cross-Variable Forecaster (f_2):** Emphasizes the interdependencies between different variables within the data, using a last-step representation to predict future values.

The OCP block's weight adjustment is depicted in part (b) of Figure 2. The long-term weights managed by EGD are designed to capitalize on the overall historical performance, while the short-term weights generated by an offline RL policy target recent changes, allowing for agile adjustment to new patterns. The combined weight $w_{t,i}$ is then applied to the forecasters' predictions to produce the final ensemble forecast.

OneNet employs a decoupled training strategy, where each forecaster is trained independently to predict future values. The OCP block is then trained to find the optimal ensemble weights that minimize the prediction error when combining the forecasters' outputs. This approach ensures that both forecasters are adequately trained even if one initially outperforms the other, a crucial feature for maintaining performance under concept drift conditions.

5.2 Model Framework

We’ve adapted OneNet to predict individual companies’ future stock prices. As discussed earlier, the most novel aspect comes from the combination of using EGD for online learning and long-term weights, and RL for offline learning and short-term weights. EGD had been traditionally used to make long-term predictions, whereas reinforcement learning has been a more recent concept. In general, RL describes a state space, in which an agent interacts with its environment and gathers rewards by learning policies in a trial and error manner. This algorithm deals with sequential decision making in a wide range of fields in both natural and social sciences, and engineering [21] but has gained more attention within financial applications. Mathematically, it can be represented as an optimization problem -

$$\max_{\pi} \mathbb{E}_{t=0}^{T-1} [R_t(s_t, a_t, s_{t+1}, x_t)] \quad (1)$$

subject to:

$$s_{t+1} = f_t(s_t, a_t, h_t) \quad (2)$$

where $a_t \in A$ indicates the actions, $s_t \in S$ the state of the system at time t , ζ_t and η_t are noise variables, and R_t is the reward received at every time step. In RL the second line in the above equation is usually referred to as the ‘environment’. The ‘agent’ intends to choose its actions a_t , given the state s_t , so as to maximize the total expected accumulated reward. Figure 7 shows this schematically.




Figure 7: Schematic representation of Environment-agent interactions

In the context of our model, at time step t , the target aims to learn short-term weights conditioned on long-term weights \mathbf{w} and experts’ performances during a short period of history $I = [l, t]$. The agent then chooses actions using a policy $\pi_{\theta_{rl}}(\mathbf{b}_t | \{w_{t,i} \tilde{y}_{i=l}^d\}_{t \in I}; \mathbf{y})$ parameterized by θ_{rl} . During training, the product between each prediction and expert weight ($w_{t,i} * \tilde{y}_i$) is concatenated with the outcome y for the input. In reference to RvS [24], the policy network is implemented as a two-layer MLP $f_{rl} : \mathbb{R}^{H \times M \times (d+1)} \rightarrow \mathbb{R}^d$, with the short-term weight calculated as:

$$\mathbf{b}_t = f_{rl}(w_{t,1}\mathbf{y}_1 \otimes \cdots \otimes w_{t,d}\tilde{\mathbf{y}}_d \otimes \mathbf{y}) \quad (3)$$

and the final ensembling weight calculated as:

$$\tilde{w}_{t,i} = (w_{t,i} + b_{t,i}) / \left(\sum_{i=1}^d (w_{t,i} + b_{t,i}) \right) \quad (4)$$

The network is then trained with the objective of minimizing the forecasting error incurred by the new weight: $\min_{\theta_{rl}} \mathcal{L}(\tilde{\mathbf{w}}) := \|\sum_{i=1}^d \tilde{w}_{t,i} f_i(\mathbf{x}) - \mathbf{y}\|^2$. As concept drift changes incrementally during inference, the expression $\mathbf{w}_{t-1} + \mathbf{b}_{t-1}$ is used to generate the prediction and train the networks after the ground truth outcome is observed:

$$\tilde{\mathbf{y}} = \tilde{w}_1 * \tilde{\mathbf{y}}_1 + \tilde{w}_2 * \tilde{\mathbf{y}}_2 \quad (5)$$

Updates to the long-term and short-term weights are as follows from the OneNet algorithm [24]:

$$w_i = w_i \exp(-\eta \|\tilde{\mathbf{y}}_i - \mathbf{y}\|^2) / \sum_{i=1}^2 w_i \exp(-\eta \|\tilde{\mathbf{y}}_i - \mathbf{y}\|^2) \quad (6)$$

$$f_{rl} \leftarrow \text{Adam}(f_{rl}, \mathcal{L}(\tilde{w}_1 * \tilde{\mathbf{y}}_1 + \tilde{w}_2 * \tilde{\mathbf{y}}_2, \mathbf{y})) \quad (7)$$

And, lastly, we update the two forecasters to be:

$$f_1 \leftarrow \text{Adam}(f_1, \mathcal{L}(\tilde{\mathbf{y}}_1, \mathbf{y})), f_2 \leftarrow \text{Adam}(f_2, \mathcal{L}(\tilde{\mathbf{y}}_2, \mathbf{y})) \quad (8)$$

Table 2: OneNet architecture

Layer	Shape	Params
Sequential-21	[-1, 64]	9,856
Linear-22	[-1, 64]	1,344
Linear-23	[-1, 64]	12,352
SiLU-24	[-1, 64]	0
SamePadConv-25	[-1, 64, 20]	0
Linear-26	[-1, 3]	195
Linear-27	[-1, 1]	65
Linear-28	[-1, 1]	65
CosineSimilarity-29	[-1]	0
ConvBlock-30	[-1, 64, 20]	0
Conv1d-31	[-1, 64, 20]	12,288
Sequential-32	[-1, 64]	9,856
Linear-33	[-1, 64]	1,344
Linear-34	[-1, 64]	12,352
SiLU-35	[-1, 64]	0
SamePadConv-36	[-1, 64, 20]	0
Linear-37	[-1, 3]	195
Linear-38	[-1, 1]	65
Linear-39	[-1, 1]	65
CosineSimilarity-40	[-1]	0
ConvBlock-41	[-1, 64, 20]	0
Conv1d-42	[-1, 64, 20]	12,288
Sequential-43	[-1, 64]	9,856
Linear-44	[-1, 64]	1,344
Linear-45	[-1, 64]	12,352
SiLU-46	[-1, 64]	0
SamePadConv-47	[-1, 64, 20]	0
Linear-48	[-1, 3]	195
Linear-49	[-1, 1]	65
Linear-50	[-1, 1]	65
CosineSimilarity-51	[-1]	0
ConvBlock-52	[-1, 64, 20]	0
Conv1d-53	[-1, 64, 20]	12,288
Sequential-54	[-1, 64]	9,856
Linear-55	[-1, 64]	1,344
Linear-56	[-1, 64]	12,352
SiLU-57	[-1, 64]	0
SamePadConv-58	[-1, 64, 20]	0
Linear-59	[-1, 3]	195
Linear-60	[-1, 1]	65
Linear-61	[-1, 1]	65
CosineSimilarity-62	[-1]	0
ConvBlock-63	[-1, 64, 20]	0
ConvNet-64	[[-1, 20, 64], [-1, 20, 64], [-1, 20, 64], [-1, 20, 64]]	0
Total		225,600
Trainable		225,600
Non-trainable		0

6 Evaluation Metric and Loss Function

As mentioned in the introduction, the Loss function we have employed is MSE loss for training the component models of RLFinNet. We have found in the literature that common evaluation metrics for general time-series forecasting models is the MSE Loss (Mean Squared Error) in the training as well as the MAE of the testing portion. For MSE, we are aiming for a mean MSE loss of all models to be lower than 0.4, which is in line with OneNet’s original implementation. While we would like to ensure that the MSE loss for all specialized models would be less than 0.4, we were not able to target this due to time constraints. Since the MAE (Mean Absolute Error) is closely linked to the MSE loss, we don’t have a specific target for the MAE, but will compare the MAE for each of our experiments. For comparable MSEs, we expect to see similar MAEs. MSE is a common loss function used in regression tasks. It measures the average squared difference between the actual values and the predicted values. For a dataset with n samples, where y_i are the actual values and \hat{y}_i are the predicted values, the MSE is calculated as:

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

The Mean Absolute Error (MAE) is another loss function used in regression tasks. It measures the average absolute difference between the actual values and the predicted values. For a dataset with n samples, the MAE is calculated as:

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

We are also plotting the predicted closing price vs. the true closing price for 3 random companies – AFGS, BA, and MDT, chosen randomly. We have also plotted the training time for each companies that we recorded in one of our experiments. All experiments were performed on a Nvidia Tesla T4 GPU on a Google Cloud virtual machine with 26GB of RAM.

7 RLFinNet Adaptations and Experiments

7.1 Data Loader Update for Finance Data

To accommodate the financial dataset described in the Dataset section, we first updated our data loader to handle multiple features including the *Open*, *High*, *Low*, *Volume*, and *Adjusted Closing* prices, along with our target forecasting variable *Close*, or closing price. The input data thus became 7-dimensional, reflecting the multivariate nature of financial time series.

7.2 Training Loop Modification

Given the extensive data across various companies, it was imperative to modify the training loop to enable sequential processing. This setup ensures that once the model completes training and making predictions for one company’s data, it automatically proceeds to the next, effectively managing multiple datasets in a single training regime. This loop not only facilitates extensive training across different datasets but also streamlines the process of validation and prediction for each company.

7.3 Implementation of Data Loader and Training Mechanism

Implementing the modified data loader involved ensuring that each company’s data is loaded, processed, and fed into the model in succession. The training mechanism was adapted to this setup, allowing for a seamless transition between training sessions for different companies without manual intervention.

261 7.4 Visualization and Post-Processing

262 Post-training, it was crucial to set up a robust framework for visualizing and post-processing the
263 results. This step included plotting the predicted vs. actual closing prices and assessing model
264 performance across different metrics such as the MSE, MAE, and prediction vs. truth plots.

265 7.5 Experimental Setup

266 With the pipeline established, we conducted a series of experiments to explore the impact of varying
267 prediction lengths and sequence lengths on the model's performance. We tested three different
268 prediction lengths: 1, 20, and 40 days, alongside three sequence lengths: 10, 20, and 40 days. These
269 experiments were designed to provide insights into the optimal configurations for both short-term
270 and long-term forecasting accuracy.

271 7.6 Insights and Observations

272 The experimental results, detailed in the subsequent section, shed light on how different configurations
273 influence the forecasting capabilities of our model. These findings are crucial for understanding the
274 trade-offs between responsiveness and accuracy in financial time series forecasting. Finding optimal
275 configurations will require many more ablations.

276 8 Results and Discussion

277 8.1 MAE for varying sequence lengths




Figure 8: MAE for each company in the dataset when trained with sequence lengths of 10, 20 and 40

278 8.2 MSE loss for varying sequence lengths




Figure 9: MSE for each company in the dataset when trained with sequence lengths of 10, 20 and 40

279 8.3 Predictions vs. Ground Truth for companies AGFS, BA, MDT




Figure 10: Predictions vs. Ground Truth with encoder sequence length=10




Figure 11: Predictions vs. Ground Truth with encoder sequence length=20




Figure 12: Predictions vs. Ground Truth with encoder sequence length=40

280 8.4 Training Time

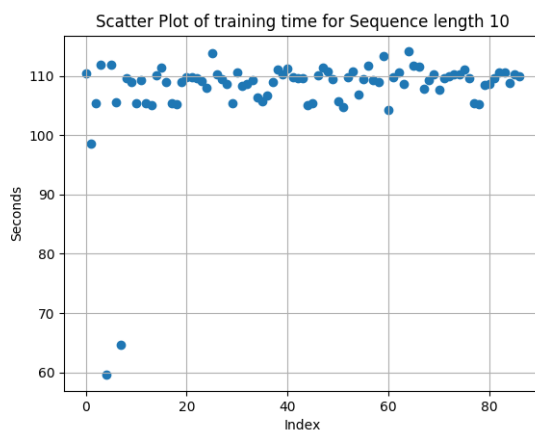


Figure 13: Training time for each specialized model

281 The original OneNet performed best with a sequence length of 40 on their hourly and monthly training
282 on the electricity transformer temperature (ETT) data, while ours did best with a sequence length of
283 20. We observe from figure 8 that the encoder sequence length of 20 had the smoothest configuration
284 for the companies, while lengths 10 and 40 exhibit some anomalies. We suspect that 10 days is
285 not enough context to make an accurate prediction, whereas 40 leads to an overfitting of the model.
286 Figure 9 demonstrates a similar pattern, where a sequence length of 20 in general had a more uniform
287 average $MSE < 0.4$, while there exists some outliers for the other two plots. This average, from out
288 literature reviews, is relatively normal within research in forecasting time-series data, though the
289 original OneNet had an MSE values around 0.35 for their ablations on the ETT dataset. Following,
290 our most insightful comparisons come from the ground truth vs. prediction plots of three random
291 companies, AGFS, BA, and MDT. All 3 companies in all three lengths were able to predict the trends
292 of the market conditions pretty closely though did not forecast the extreme variances that exist in
293 the stock market. This is expected since our training tries to model the distribution of the data and
294 cannot precisely predict when spikes occur in the sequence of data. Since there is scarce research
295 done on forecasting specifically stock price data using MSE or MAE as metrics, we consider our
296 model relatively new within the domain and omit comparisons with previous baseline models. To
297 assess the computational efficiency of RLFinNet, we plot the training times with encoder sequence
298 length 10 as a reference in Figure 13. and observe that each company less than 2 minutes to train on,
299 making our model quick and lightweight.

300 9 Conclusion and Future Works

301 From our experimentation, we have observed that RLFinNet effectively deals with concept drifts and
302 visually performs relatively well in forecasting performances of stock prices. At only 225k parameters
303 with our current implementation, it is very lightweight easily accessible for each company. Our
304 prediction vs. ground truth plots demonstrate that the predictions are able to model the volatility of
305 the stock market but not to the extreme variances, but this is to be expected due to the non-stationary
306 nature of the financial market. While RLFinNet is designed to train on real-time data, we have only
307 worked with public stock price datasets for 87 different companies. We aim to configure our network
308 on a real trading floor and streamline financial data from the online market through various APIs.
309 Additionally, we hope to automate the process of hyper-parameter tuning to be tailored for trading
310 users in their companies of interest.

References

- [1] Ahmadi, S. (2024). A Comprehensive Study on Integration of Big Data and AI in Financial Industry and its Effect on Present and Future Opportunities. *International Journal of Current Science Research and Review*, 07 (01), 66-74.
- [2] Alemansour, M.M., & Rastgarpour, M. (2021). Improving the Prediction of Cryptocurrencies Movement Using Adversarial Training. In M. Rastgarpour (Ed.), *Progress in Intelligent Decision Science: Proceeding of IDS 2020* (pp. 282-293). Springer International Publishing.
- [3] Azhikodan, A.R., Bhat, A.G.K., & Jadhav, M.V. (2019). Stock trading bot using deep reinforcement learning. In *Innovations in Computer Science and Engineering: Proceedings of the Fifth ICICSE 2017* (pp. 41-49). Springer Singapore.
- [4] Cannelli, Loris, Giuseppe Nuti, Marzio Sala, and Oleg Szehr. "Hedging using reinforcement learning:331 Contextual k-armed bandit versus Q-learning." *The Journal of Finance and Data Science* 9 (2023): 10010
- [5] Chaouki, A., Hardiman, S., Schmidt, C., Sérié, E., & De Lataillade, J. (2020). Deep deterministic portfolio optimization. *The Journal of Finance and Data Science*, 6 (1), 16-30.
- [6] Chen, Y.-Y., Chen, C.-T., Sang, C.-Y., Yang, Y.-C., & Huang, S.-H. (2021). Adversarial attacks against reinforcement learning-based portfolio management strategy. *IEEE Access*, 9, 50667-50685.
- [7] Cheng, L.-C., Huang, Y.-H., & Wu, M.-E. (2018). Applied attention-based LSTM neural networks in stock prediction. In *2018 IEEE International Conference on Big Data (Big Data)* (pp. 4716-4718). doi: 10.1109/BigData.2018.8622541.
- [8] Emmons, S., Eysenbach, B., Kostrikov, I., & Levine, S. (2022). Rvs: What is essential for offline rl via supervised learning? *ICLR*.
- [9] Feng, F., Chen, H., He, X., Ding, J., Sun, M., & Chua, T.-S. (2019). Enhancing Stock Movement Prediction with Adversarial Training. In *IJCAI*, 19, 5843-5849.
- [10] Jia, W. U., Chen, W. A. N. G., Xiong, L., & Hongyong, S. U. N. (2019). Quantitative trading on stock market based on deep reinforcement learning. In *2019 International Joint Conference on Neural Networks (IJCNN)* (pp. 1-8). IEEE.
- [11] Jiang, W. (2021). Applications of deep learning in stock market prediction: recent progress. *Expert Systems with Applications*, 184, 115537.
- [12] Jiang, Z., Xu, D., & Liang, J. (2017). A deep reinforcement learning framework for the financial portfolio management problem. *arXiv preprint arXiv:1706.10059*.
- [13] Liang, Z., Chen, H., Zhu, J., Jiang, K., & Li, Y. (2018). Adversarial deep reinforcement learning in portfolio management. *arXiv preprint arXiv:1808.09940*.
- [14] Liu, X.-Y., Yang, H., Chen, Q., Zhang, R., Yang, L., Xiao, B., & Wang, C.D. (2020). FinRL: A deep reinforcement learning library for automated stock trading in quantitative finance. *arXiv preprint arXiv:2011.09607*.
- [15] Liu, X.-Y., Xiong, Z., Zhong, S., Yang, H., & Walid, A. (2018). Practical deep reinforcement learning approach for stock trading. *arXiv preprint arXiv:1811.07522*.
- [16] Moody, J., & Saffell, M. (2001). Learning to trade via direct reinforcement. *IEEE transactions on neural Networks*, 12 (4), 875-889.
- [17] Meng, T.L., & Khushi, M. (2019). Reinforcement learning in financial markets. *Data*, 4 (3), 110.
- [18] Neuneier, R. (1995). Optimal asset allocation using adaptive dynamic programming. *Advances in neural information processing systems* 8.
- [19] Scott, E., Eysenbach, B., Kostrikov, I., & Levine, S. (2022). Rvs: What is essential for offline rl via supervised learning? *ICLR*.
- [20] Shaban, W.M., Ashraf, E., & Slama, A.E. (2024). SMP-DL: a novel stock market prediction approach based on deep learning for effective trend forecasting. *Neural Comput & Applic*, 36, 1849–1873.

- 359 [21] Sutton, R.S., & Barto, A.G. (1999). Reinforcement learning: An introduction. *Robotica*, 17 (2),
360 229-235.
- 361 [22] World Bank. "Market Capitalization of Listen Companies." Retrieved from <https://data.worldbank.org/indicator/CM.MKT.LCAP.CD/>.
- 363 [23] Xu, Y., & Cohen, S. (2018). Stock movement prediction from tweets and historical prices. In
364 *ACL*, 1, 1970–1979.
- 365 [24] Wen, Q., Chen, W., Sun, L., Zhang, Z., Wang, L., Jin, R., & Tan, T. (2024). Onenet: Enhancing
366 time series forecasting models under concept drift by online ensembling. *Advances in Neural*
367 *Information Processing Systems* 36.
- 368 **Work Division:** ¹ **GitHub Repository:** ²

¹Naman Choudhary: Research OneNet + RLFinNet adaption, Karthik Subramaniam: Research OneNet + RLFinNet ablations/adaption, Nimish Jindal: Research OneNet/Adv-ALSTM baseline + RLFinNet adaption, Sarah Li: Research OneNet + RLFinNet adaption

²GitHub